

Synchronisation Properties

Mutual Exclusion: if Pi in CS, all processes cannot enter CS
Progress: if none in CS, one waiting enters CS

Bounded Wait: upper bound on waiting time

Independence: Process not in CS should not block other processes

Incorrect Synchronisation

Deadlock: all processes blocked

Livelock: processes running, but

changing state to avoid deadline, doing useless

Starvation: wait forever

TestAnd Set Reg, MemLocat

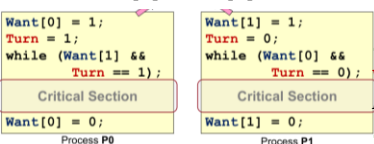
Atomically:

Load [MemLoc] into [Reg]

[MemLoc] <- 1

Peterson's Algo

- Busy Waiting
- Need variable setting
- Only 2 processes
- Turn=Want[0]=Want[1]=0



Semaphore

sem_t a;
sem_init(&a, process=1, val=1)
(process=1, thread=0)
sem_wait(&a); (--, else block)
sem_post(&a); (++)/signal)

Scurr = Sinit + #signal - #wait

Producer Consumer

Producer produces if buffer < K
Consumers eat if > 0

count=in=out=0

mutex=S(1), notFull=S(K)

notEmpty=S(0)

Producer

wait(notFull)

wait(mutex)

CS

signal(mutex)

signal(notEmpty)

Consumer

wait(notEmpty)

wait(mutex)

CS

signal(mutex)

signal(notFull)

Reader Writer

Reader must have X lock

Writer must have S lock

all mutexes = S(1)

Reader-preference

writer()

wait(roomEmpty);

Modifies data

signal(roomEmpty);

reader()

wait(mutex);

nReader++;

if (nReader == 1)

wait(roomEmpty);

signal(mutex);

Reads data

wait(mutex);

nReader--;

if (nReader == 0)

signal(roomEmpty);

signal(mutex);

Writer-preference

writer()

wait(writeMut);

writeCount++

if (nWriter == 1)

wait(readTry);

signal(writeMut);

wait(roomEmpty);

Modifies data

signal(roomEmpty);

wait(writeMut);

nWriter--;

if (nWriter==0)

signal(readTry);

signal(writeMut);

reader()

wait(readTry);

wait(readMut);

nReader++;

if (nReader == 1)

wait(readMut);

signal(readMut);

signal(readTry);

Reads data

wait(readMut);

nReader--;

if (nReader == 0)

signal(roomEmpty);

signal(readMut);

Thread Mutex

pthread_mutex

pthread_mutex_lock()

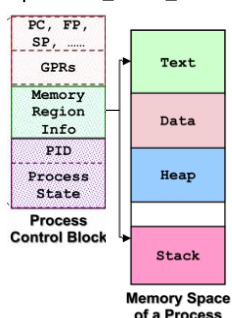
pthread_mutex_unlock()

pthread_cond

pthread_cond_wait()

pthread_cond_signal()

pthread_cond_broadcast()



Memory Usage

Transient Data eg params, local vars
Persistent Data eg global vars

OS Handles:

- Allocate memory space to new processes
- Manage memory space for processes
- Protect memory space of processes from each other
- Provides memory related system calls to processes
- Manage memory space for internal use

Main Topics:

Memory Abstraction

Presenting a logical interface for memory accesses

Contiguous Memory Allocation

Allocating and managing continuous chunk of memory

Disjoint Memory Allocation

Allocating and managing memory in disjoint areas

Virtual Memory Management

Use of secondary storage as an extended memory region

Memory Abstraction

- Cannot fix addr during comp time

- Need to protect address space when

2 processes have same phys mem

Address Relocation (not used)

- diff to distinguish mem ref from int ref

Base + Limit Registers (not used)

- Base reg added to all refs

- Limit reg needed for mem integrity

- Have to add and compare every inst

Multitasking

- Allow multiple processes at same time to allow switching

- Free up memory by removing terminated processes, swap blocked processes to sec storage

Contiguous Mem Mgmt

- Contiguous process space, phys mem can store whole memory

Fixed-Size Partitions

- Physical memory is split into fixed number of partitions of equal size

- A process will occupy one of the partitions

- Internal Fragmentation

Variable-Size Partitions

- Partition is created based on the actual size of process

- OS keep track of the occupied and free memory regions

- Perform splitting and merging when necessary (no internal partition)

- Need to store more info

- Need to locate regions

Dyn Algorithms: FF, BF, WF, NF

Buddy Algorithm

- Free blocks split in half to find smallest block of $2^S \geq N$

- Buddies have 0..S-1 bits equal

- kth level buddy has kth bit flipped

- $32B = 2^5$ root is 00000, 10000 (16B buddies)

FF, BF, WF, NF: EF, no IF

Fixed: IF, no EF

Buddy: IF and EF

Translation Lookaside Buffer (TLB)

Caches for Page Table Entries (flush when switching processes)

Average memory access time (AMAT):

= TLB hit + TLB miss = $90\% \times (1ns + 50ns) + 10\% \times (1ns + 50ns + 50ns)$

= $P(TLB \text{ hit}) \times \text{latency}(TLB \text{ hit}) + P(TLB \text{ miss}) \times \text{latency}(TLB \text{ miss}) = 56ns$

Disjoint Mem with Paging

- Logical Memory split into logical pages of same size

- Physical memory split into physical frames of same size

- During runtime, pages are loaded into available memory frames

Paging Scheme

- needs to know starting address

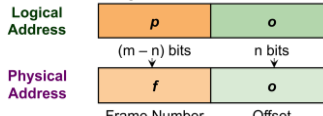
and size of process

- EF not possible. IF insignificant

Phys frame size == logical page size == 2^k

$PA = F \times \text{sizeof}(\text{physical_frame}) +$

offset



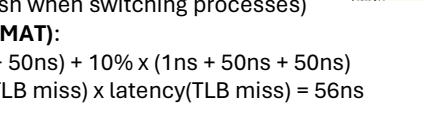
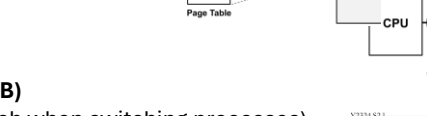
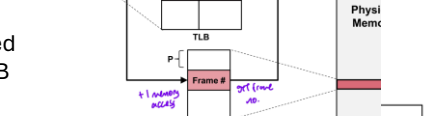
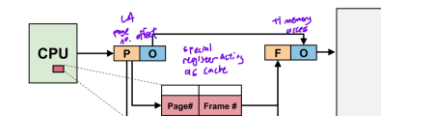
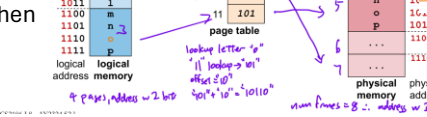
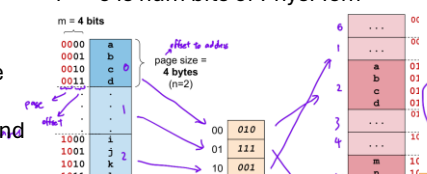
p finds which page in page table

f is what's inside page table

o is num bits of a page (1 page is 2^B)

$p \times o$ is num bits of LogMem

$f \times o$ is num bits of PhysMem



Paging Scheme Protection

Access-Right Bits – RWX. Every memory access checks (not every process uses all the pages)

Valid Bit – whether is valid to access by process

Page Sharing

Several processes can share the same physical memory frame

Copy-on-Write

Segmentation

Some regions may grow / shrink at execution time

Each memory segment is specified with **SegmentName + Offset**, and specifies a limit

LA <SegID, Offset>:

SegID looks up <Base, Limit> in segment table

PA = Base + Offset

Offset < limit

