


## Monolithic Architecture




## Microkernel (isolated, but low perf)

Provides IPC, Addr space/Thread mgm

## High Level Services


run as server process outside OS, communicate using IPC, Better isolation between kernel and HLS





## Virtual Machine / Hypervisor

Virtualization of underlying hardware

### Type 1 (individual VM to guest OS)



### Type 2 (run in host OS, guest run in VM)



## Instruction Execution

Memory indicated by PC Instruction, read from mem or GPR

### Memory Context:

- text (instructions),
- data (global vars)
- Stack, Heap

### Hardware Context:

GPR, PC, SP, FP

### OS Context:

PID, Process State

### Stack Memory

Stores function invoc in a **stack frame**

Contains:

- Return addr of caller
- function args
- local vars

### Saved Registers


(register spilling)

- Use memory to hold GPR

- GPR can be used for other purposes and

restored later

### Stack Frame



### Exponential Average

Predicted<sub>n+1</sub> = αActual<sub>n</sub> + (1-α)Predicted<sub>n</sub>

- Actual<sub>n</sub> = The most recent CPU time consumed

- Predicted<sub>n</sub> = past history

- α = Weight placed on recent event

- Predicted<sub>n+1</sub> = Latest prediction

## Stack Frame

### Setup/Teardown

On executing function call:

- Caller: Pass arguments with registers and/or stack
- Caller: Save Return PC on stack

Transfer control from caller to callee

- Callee: Save registers used by callee. Save old FP, SP

- Callee: Allocate space for local variables of callee on stack

- Callee: Adjust SP to point to new stack top

On returning from function call:

- Callee: Restore saved registers, FP, SP

- Transfer control from callee to caller using saved PC

- Caller: Continues execution in caller

### Dynamic Memory (Heap)

Runtime known only at compile time (not data), no definite allocation/dealloc timing (not stack)

### PCB

Maintained by Kernel

PCB<sub>1</sub>, PCB<sub>2</sub>, PCB<sub>3</sub>, ...

PC, FP, SP, ...

GPRs

Memory Region Info

PID

Process State

Process Control Block

Process Table

Text

Data

Heap

Stack

Memory Space of a Process

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

## Syscalls

1. Library places syscall no. in a register

2. Switch from user to kernel

3. Dispatcher uses syscall no. as index to execute handler

4. Return control to library

5. Switch from kernel to user

### Exceptions

- Synchronous

- Execute exception handler

### Interrupt

- Asynchronous,

- Usu. hardware related

- Suspend program execution

- Executes interrupt handler

### int fork()

Returns child pid for parent,

Returns 0 for child

int main(int argc, char\* argv[]);

argc: num of args, inc program name

argv: each element is a C char string

- If 3 params, argc = 4

- arg0 is a.out, arg1-3 are params

int execl(const char \*path,

const char \*arg0, ..., NULL);

execl("/bin/l", "l", "-l", NULL);

Path = "/bin/l", arg0 = "l",

arg1 = "-l"

void exit(int status);

0 = normal, !0 = problem

int wait(int \*status);

blocks then returns PID of terminated child

returns immediately if no child

Criteria: Balance, Fairness (starve)

### Response Time:

first CPU time – arrival time

Turnaround Time: finish – arrival

Waiting Time: time waiting for CPU

Throughput: Number of tasks finished per unit time

CPU Utilisation: %time CPU is used

Waiting Time: turnaround – CPU work done – IO waiting time

...

...

...

...

...

...

...

...

...

...

## Zombie Process

Child that exited but parent not called wait() by parent

### If Parent process terminates before child process:

- init process becomes "pseudo" parent of child processes

- Child termination sends signal to init, which utilizes wait() to cleanup

### Child process terminates before parent but parent did not call wait:

- Child becomes zombie

- Can fill up process table

- May need a reboot to clear the table

### Processing Environment

1. Batch:

- No interaction required, No need to be responsive

2. Interactive (or Multiprogramming):

- With active user interacting with system

- Should be responsive, consistent in response time

3. Real time processing:

- Have deadline to meet

- Usually periodic process

### Lottery

a process holding X% of tickets uses resource X%

Responsive:

- A newly created process can participate in the next lottery

- Provides good level of control:

- A process can be given Y lottery tickets depending on priority, and to be distributed to child,

- Can control the proportion of usage

- Each resource can have its own set of tickets

- Different proportion of usage per resource per task

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

## Batch

FCFS: no starvation

Convoy Effect: all wait CPU, all wait IO, repeat (a lot of idling)

### Shortest Job First:

Minimises ave waiting time

Need to know CPU time

Long jobs starve

Uses exponential average

### Shortest Remaining Time (Preemptive)

Good service for late short jobs

Only really need to check when new job comes

### Interactive

- Interval of Timer Interrupt (ITI), invoke scheduler every x ms

- Time Quantum (a multiple of ITI), execution duration

NB Context switching wastes time

### Round Robin (RR)

Store in FIFO until TQ, give up or block

Placed to back of queue

Time before a task gets CPU bounded by (n - 1)q

Big quantum: Better CPU utilization but longer waiting time

Small quantum: Bigger overhead (worse CPU utilization) but shorter waiting time

### Priority Scheduling

Preemptive: higher preempts low

Non: latecomer waits for next round

Low priority can starve

Priority Inversion: Resource locked, lower priority preempts higher priority

### Multilevel Feedback Queue (MLFQ)

- Response time for IO bound processes minimised

- Turnaround time for CPU bound processes minimised

- If P(A)>P(B), A runs

- If P(A)=P(B), A and B runs in RR

- Priority reduced if fully utilizes time slice, else retain Priority

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

**Multithreaded processes**

- Shared Memory Context: Text, Data, Heap
- Shared OS Context: Process id, stack
- Unique information needed for each thread**

- thread id
- Registers (GPR, special)

**Process Context Switch:**

OS, hardware, memory context

**Process Thread switch:**

Hardware context: registers, FP, SP

**Benefits**

Share resources

Less kernel mode switches

**User Thread:**

- Thread is implemented as a user library
- A runtime system (in the process) will handle thread related operation
- Kernel is not aware of the threads in the process – 1 kernel thread to 1 process entryptpoint even if more than 1 process thread

**Kernel Thread:**

- Thread implemented in OS
- Thread operation is handled as system calls
- Thread-level scheduling is possible:
- Kernel schedule by threads, instead of by process
- Kernel may make use of threads for its own execution

**User Thread Adv and Disadv**

- + Multithreaded program on any OS
- + more configurable, flexible
- OS not aware of threads, scheduling performed at process level
- One thread blocked = all threads blocks

**Simultaneous Multithreading (SMT)**

Multiple sets of registers to allow threads to run in parallel

**Create pthread**

```
int pthread_create(
    pthread_t* tidCreated,
    const pthread_attr_t* threadAttributes,
    void* (*startRoutine) (void*),
    void* argForStartRoutine );
returns 0=success, !=0=error
```

**Exit pthread**

```
int pthread_exit( void* exitValue );
if return XYZ instead, then use XYZ
instead of pthread_exit
```

**Join pthread (synchronise)**

```
int pthread_join( pthread_t threadID,
void **status );
```

**Shared Memory**

**Advantages:**

- Only create and Attach shared memory region) involves OS
- Ease of use:
- Shared memory region behaves the same as normal memory space
- i.e. Information of any type or size can be written easily

**Disadvantages:**

- Synchronization:
- Shared resource -> Need to synchronize access
- Implementation is usually harder

**Usage**

1. Create/locate a shared memory region  
`M shmId = shmget( IPC_PRIVATE, 40, IPC_CREAT | 0600);`
  2. Attach M to process memory space  
`shm = (int*) shmat( shmId, NULL, 0 )`
  3. Read from/Write to M (Values written visible to all process that share M)
  4. Detach from memory space after use  
`shmdt( (char*) shm);`
  5. Destroy `shmctl( shmId, IPC_RMID, 0);`
- Only one process need to do this  
Can only destroy if M is not attached to any process

**Direct Communication**

Send(p2, msg)  
Receive(P1, msg)  
Need to know pid to send (which is not fixed). Instead, send to MailBox, which can be shared among processes

**Blocking Primitives (Synchronous):**

- Send(): Sender is blocked until the message is received
- Receive(): Receiver is blocked until a message has arrived

**Non-Blocking Primitives (async):**

- Send(): Sender resume operation immediately
- Receiver(): Receiver either receive the message if available or some indication that message is not ready yet

**Advantages:**

- Portable: easily implemented on different processing environments
- Easier sync when synchronous primitive is used

**Disadvantages:**

- Inefficient: usu require OS
- Harder to use: limit msg size and format

**Pipes**

Pipe functions as circular bounded byte buffer with implicit synchronization:

- Writers wait when buffer is full
- Readers wait when buffer is empty

**int pipe( int fd[] )**

fd[0]: reading end, fd[1]: writing end

**Pipe Code**

```
int pipeFd[2], pid, len;
char buf[100], *str = "Hello There!";
pipe( pipeFd );
if ((pid = fork()) > 0) { /* parent */
    close(pipeFd[READ_END]);
    write(pipeFd[WRITE_END], str,
    strlen(str)+1);
    close(pipeFd[WRITE_END]);
} else { /* child */
    close(pipeFd[WRITE_END]);
    len = read(pipeFd[READ_END], buf,
    sizeof buf);
    printf("Proc %d read: %s\n", pid, buf);
    close(pipeFd[READ_END]);
}
```

**Signals**

A form of inter-process communication

- An asynchronous notification regarding an event
- Sent to a process/thread
- The recipient of the signal must handle the signal by:
- A default set of handlers OR
- User supplied handler (for some only)

**Signal Code**

```
void myOwnHandler( int signo ) {
    if (signo == SIGSEGV){
        printf("Memory access blows up!\n");
        exit(1); }
}

int main(){
    int *ip = NULL;
    if (signal(SIGSEGV, myOwnHandler) ==
    SIG_ERR)
        printf("Failed to register handler\n");
    *ip = 123;
    return 0;
}
```

**Input Redirection**

```
int fp_in = open("./file.txt",
O_RDONLY);
int fp_out = open("./talk.out",
O_CREAT | O_WRONLY);

if(fork() == 0) {
    dup2(fp_in, STDIN_FILENO);
    dup2(fp_out, STDOUT_FILENO);
    execlp("./talk", "talk", (char *) 0);
    close(fp_in);
    close(fp_out);
}
else
    wait(NULL);
```

**To simply redirect output from**

**stdout:**  
`close(STDOUT);`  
`open("result.txt", O_WRONLY |`  
`O_CREAT | O_TRUNC, 0644);`  
`printf("hello world") // will be written`  
`into result.txt file`

**PIPE**

```
int p[2];
char str[] = "Hello this is the parent.";
// This creates a pipe. p[0] is the
reading end,
// p[1] is the writing end.
if(pipe(p) < 0)
    perror("lab2p2e: ");
// We will send a message from father
to child
if(fork() != 0) {
    // Parent
    close(p[0]); // The the end we are not
    using.
    write(p[1], str, strlen(str));
    close(p[1]);
    wait(NULL);
}
else
{
    // Child
    char buffer[128];
    close(p[1]); // Close the writing end
    read(p[0], buffer, 127);
    printf("Child got the message
    \"%s\"\n", buffer);
    close(p[0]);
}
```