

<p>Latency</p> <p>IAAS, PAAS, SAAS</p> <p>Cloud-native Fts</p> <ul style="list-style-type: none"> - Immutable infra - Microsvcs - API driven - Service mesh - Containers - Dynamically mgmt <p>Development</p> <p>CI/CD</p> <p>DevOps</p> <p>Serverless</p> <p>Stack</p> <ul style="list-style-type: none"> - Infra - Provisioning - Runtime - Orchestration, mgmt - App defn, devt - Observability <p>Deployment</p> <ul style="list-style-type: none"> - Integration - Content delivery - Heterogeneous plats - Coordination - Security <p>Containers</p> <ul style="list-style-type: none"> - Runtime with code (interop, portable) - Dep, change mgmt - Env mgmt - Reproducible - Isolation, security - Orchestratable <p>Orchestrator (K8s)</p> <ul style="list-style-type: none"> - Provisioning, deployment (coordn, fault tol) - networking - scaling - availability - lifecycle mgmt <p>Serverless (AWS λ)</p> <ul style="list-style-type: none"> - via containers - Func-AAS (FAAS) - Event-driven execution model - Stateless, ephemeral, managed by PAAS <p>SDLC</p> <ul style="list-style-type: none"> - Waterfall - Agile 	<p>CI/CDel/Dep</p> <p>Del: manual dep</p> <p>Dep: auto dep</p> <ul style="list-style-type: none"> - Low risk release - Faster to mkt - Higher qual - Reliable - Scalable - Lower cost <p>Devops</p> <ul style="list-style-type: none"> - Commit, build, test, release artifacts, - CI/CD/monitoring/communication <p>Requirements</p> <p>Elicit, Analyse, Specify, Validate</p> <p>SWE Req Spec</p> <ol style="list-style-type: none"> FRs NFRs System Req Constraints External Interfaces <p>SRS Qualities</p> <ol style="list-style-type: none"> Accurate Complete Modifiable Ranked Testable Traceable Unambiguous Verifiable Valid <p>Business Req</p> <p>E.g. Reduce costs by 25%</p> <p>User Req</p> <ul style="list-style-type: none"> - As a (user), I can (action) so that I can (action) - [ref] The system supports png and-jpg <p>FR: what system does</p> <p>NFR: how well works (reliability, efficiency, usability)</p> <p>System Reqs</p> <p>Hardware Systems</p>	<p>External Interfaces</p> <p>Connections to outside world</p> <p>Constraints</p> <p>e.g. Backward compatible, mem usage</p> <p>External Quality Attr</p> <ol style="list-style-type: none"> Availability Safety Deployability Usability Performance Security Compatibility Reliability Installability Integrity Robustness Interoperability <p>Internal Quality Attr</p> <p>Felt by devs, maintainers</p> <ol style="list-style-type: none"> Testability Efficiency Modifiability Scalability Portability Reusability Verifiability Maintainability <p>Security</p> <ul style="list-style-type: none"> - privacy - authentication - integrity <p>Performance</p> <ul style="list-style-type: none"> - Response time - Throughput - Data capacity - Concurrent users - Predictability - Latency - Behaviour when overloaded <p>Scalability</p> <ul style="list-style-type: none"> - hard/software - Vert scaling (bigger) - Hori scaling (more) <p>Availability</p> <ul style="list-style-type: none"> - up time/total - Five 9s=99.999% - needs redundancy 	<p>Usability</p> <ul style="list-style-type: none"> - UX - num clicks to do X - ease of learning/use <p>Traceability</p> <ul style="list-style-type: none"> - type, ID - descr - priority - etc <p>Validation</p> <ul style="list-style-type: none"> - biz-relevance <p>Verification</p> <ul style="list-style-type: none"> - requirements are correct (complete, correct, feasible, prioritised) <p>Control Flow</p> <p>Data Flow</p> <p>Call and Return</p> <ul style="list-style-type: none"> - movement of control <p>Message</p> <ul style="list-style-type: none"> - Sync <p>Event</p> <ul style="list-style-type: none"> - emitted for listeners - immutable - ordered <p>Architecture Decomposition</p> <p>Hori slicing by layers</p> <p>Vert slicing by fts</p> <p>Cohesion Types</p> <p>Functional, Layer, Comms, Seq, Procedural, Temporal, Utility</p> <p>Coupling Types</p> <p>content, common, control, data, external, temporal, import</p> <p>Technical Partitioning</p> <ul style="list-style-type: none"> - SoC - Aligns with expertise <p>Domain Partitioning</p> <ul style="list-style-type: none"> - align with domain components map to problem 	<p>Layered Architecture</p> <ul style="list-style-type: none"> - monolith + technical partitioning - n-tier architecture with different layers, (+) perf optim, min context switching (-) less modifiable, new function changes multiple parts <p>Modular Monolith</p> <ul style="list-style-type: none"> - monolith + domain partitioning <p>Event-Driven</p> <ul style="list-style-type: none"> - Distributed + Technical partitioning <p>Microservices</p> <ul style="list-style-type: none"> - Distributed + Domain Partitioning <p>Pipe and Filter</p> <ul style="list-style-type: none"> - Data enters and flows through components as a <u>stream</u> (incremental computation). - Independent filters - Data source → filter → data sink (via pipes) - Useful for batch processing <p>MVC</p> <ul style="list-style-type: none"> - V: UI (http response) - C: coords M and V (handles http, select model, preps view) - M: business logic + persistence (server) V → C: user actions C → V: update model M → V: M triggers V update V → M: query for state <p>Benefits of MVC</p> <ul style="list-style-type: none"> - SoC: modularity btwn app state and processing - Extensibility: new V/C pair can be added - Reduced comm complexity - Testable/mockable <p>Single Page App (SPA)</p> <ul style="list-style-type: none"> - JS program in browser w/o refreshing - less perceived latency 	 <p>Architecture Diags</p> <ul style="list-style-type: none"> - big picture map Software system <u>Container</u> is the context or boundary which executes code or stores data Components group related functionality behind interface <p>REST</p> <ul style="list-style-type: none"> - REpresentational State Transfer - stateless - self contained - cacheable for network efficiency - application is layered - Uniform API interface -- Via resource identifiers <p>Pros and Cons</p> <ul style="list-style-type: none"> (+) less tightly coupled (+) scalable, usable, accessible, mixable (-) repeated queries (-) stale cache (-) unspecific data (cf GraphQL) <p>DDD</p> <p>Interactions between contexts</p> <p>model the interactions in the subdomains</p> <p>Domain: the problem space</p> <p>Sub-domain: component of main domain, focusing on specific set of resp; reflect biz org struct</p>	<p>Bounded Context:</p> <p>Highly cohesive boundary, includes input, output, events, reqs, procs, data models relevant to sub domain.</p> <p>Aggregate</p> <ul style="list-style-type: none"> - Cluster of related objs treated as a single obj for data changes - High cohesion, derived from business reqs - Xact boundary: all changes succeed or none - Consistency boundary: externals read state, or execute public interface methods <p>Aggregate Root</p> <ul style="list-style-type: none"> - parent entity of cluster <p>DDD Shared Kernel Collaboration</p> <p>Two contexts devt independently but end up overlapping</p> <p>DDD Upstream-Downstream Collab</p> <p>(provider-consumer r/s)</p> <p>DDD Comms</p> <p>Integration between BC are sync or async</p> <p>Event Storming</p> <ul style="list-style-type: none"> - Domain events are simply some sort of event that occurs in the domain & that are relevant. - A command causes an event - Find aggregates - Policy (WHEN event THEN command) Cmd-Ag-Ev → policy → Cmd-Ag-Ev E.g. when the payment card is submitted, then confirm order. - Draw boundaries <p>Database-server-per-Service</p> <ul style="list-style-type: none"> + loose coupling, scaling at db level + easier to replace underlying DB - expensive <p>Data Independence: private schema/tables</p>	<p>enforce data security at db levels e.g. db users</p> <p>Data Delegate Pattern</p> <p>service which all others look to for a db access</p> <p>Data Lake pattern</p> <p>read-only, queryable data sink. Shared space containing (copy of) data from concerned microsvcs</p> <p>Saga (Data) Pattern</p> <ul style="list-style-type: none"> - Every action has a compensatory action - If error on stage 4, compensate earlier stages - Note not ACID - Sequence matters - Hard to compensate actions (eg notifs) to end <p>Event Sourcing (Data)</p> <ul style="list-style-type: none"> - Events: Source of Truth - System state: projected from series (use snapshots for perf) - append mutations to Event Store (ID, eventType, data) - notify multiple subscribers to build projections - for scalability, time travel, auditability, traceability <p>Command Query Resp Segregation (CQRS)</p> <p>Cmd mutations return nth Queries read, return value</p> <ul style="list-style-type: none"> - Separate read/write models (imbal R/W) - E.g. write to Kafka event stream via KSQL, materialise as precomputed query to update (multiple) DB - Eg read model is adiff table or denorm Mongo w push/pull from EventStore <p>Service Instance per Host/VM/Container</p> <p>Devops Practices</p> <ul style="list-style-type: none"> - Local workstations to hosted infra - Provision env, deploy to hosted - Cheap for new env 	<p>Immutable Infra</p> <p>All components reprovision and recreate</p> <p>Infra as Code</p> <p>Orchestration: K8s</p> <p>Config: Ansible/Terraform</p> <p>Svc Collab: Orchestration vs Choreography</p> <p>central brain vs inform each its job</p> <p>Svc Comm:</p> <p>Sync/Async IPC Req/Response</p> <ul style="list-style-type: none"> - sync/async <p>Notif: (one way)</p> <p>Event-driven</p> <p>Real time data streams, loose coupling</p> <p>API Gateway</p> <ul style="list-style-type: none"> - single entry pt into system - tailored API - other responsibs like authenticates, monitoring, LB, caching <p>Backend for Frontend (BFF)</p> <ul style="list-style-type: none"> - diff gateway for diff platform (web, mobile) <p>Svc Discovery</p> <p>API Gateway needs to know IP addr, port of each microSvc.</p> <p>Client-Side Disc</p> <ul style="list-style-type: none"> - Svcs register w svc registry - Queries registry - LB algo to select <p>Svr-Side Disc</p> <ul style="list-style-type: none"> - Svcs register w svc registry - Client requests svc via router/LB. - LB queries svc reg, routes requests to service 	 <p>Svc Registry Pattern</p> <ul style="list-style-type: none"> - Svc Registry knows instances, addr, port. - Instances register on startup, dereg on shutdown <p>Initiating Event:</p> <p>from end user which kicks off business process</p> <p>Derived Event:</p> <p>internal s/w event triggered by init event</p> <p>Event Structure</p> <ul style="list-style-type: none"> - Unkeyed: stmt of fact - Entity: K-V pair - Keyed: not an entity, used for partitioning data stream for data locality within a single partition <p>Async Comm</p> <ul style="list-style-type: none"> - more responsive - more available - more complex error handling <p>Event-Driven Arch</p> <ul style="list-style-type: none"> - Producer (pub), Broker, Consumer (sub) - Bounded Contexts impt for microsvcs, but nice-to-have - Built on event processing i.e. responding to sth that happened, and trigger more events (Microservices based on responding to requests) <p>Event Broker</p> <ul style="list-style-type: none"> - recvs events, holds in queue/partition, sends to consumer - immutable, ordered append-only
--	--	---	--	--	---	--	--	--	--

