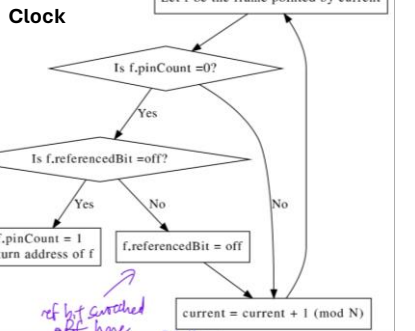




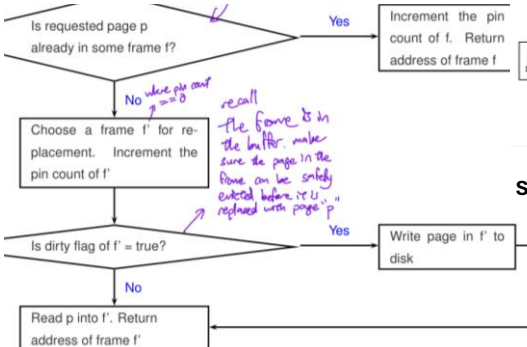
DBMS stores data on non-volatile disk for persistence
 DBMS processes data in main memory (RAM)
 read: transfer data from disk to RAM
 write: transfer data from RAM to disk

Disk access time:

- seek: moving arms to position disk head on track
- rotational delay: waiting for block to rotate under head
- =time for one rev / 2
- transfer: moving data to/from disk surface
- =num sectors * (time for one rev / num sectors per track)
- access = **seek + rotational delay + transfer**



Slotted Page (un)/packed




Storing records

Heap file: unordered file
 Sorted file: ordered on search key
 Hashed file: hash function

Heap file Implementation

LL vs Page dir (below 2 diagrams)



B+ Tree

Order of index d
 Root node has [1,2d] entries (up to 2d+1 pointers)
 Nonroot has [d,2d] entries (up to 2d+1 pointers)

Format

Format 1: k* is an actual record (clustered)
 Format 2: k* is of the form (k, rid)
 Format 3: k* is of the form (k, rid-list)

Split Overflow

- Split overflowed leaf node by distributing d+1 entries to new leaf node (i.e. right node should have 1 more records)
- Create a new index entry using the smallest key in new leaf node
- Insert new index entry into parent node of overflowed node

Split Insert leaf

First d into N, last d+1 into node to right (push middle+ entry to new node)
 Then insert internal

Insert Internal (pull middle up)

If no parent, create parent.
 Pull middle up, middle+ entry goes right

Redistribute leaf (overflow)

If non-full adj right, first 2d goes in, last entry goes right
 If non-full adj left, last 2d goes in, first entry goes left
 Else split

Merge Delete w dist (pull parent down)


Move first entry from right
 If Move last entry from left
 Else move entries to right
 Else move to left

Merge underflow

Merge right
 If merge left
 Else

Bulk Load

1. Sort the data entries to be inserted by search key
2. Load the leaf pages of B+-tree with sorted entries
3. Initialize B+-tree with an empty root page
4. For each leaf page (in sequential order), insert its index entry into the rightmost parent-of-leaf level page of B+-tree



- **GetBucketNum(k)** returns bucket # where entry with search key k is located

$$GetBucketNum(k) = \begin{cases} h_{level}(k) & \text{if } h_{level}(k) \geq next, \text{ \& bucket not split, or else ok this} \\ h_{level+1}(k) & \text{otherwise} \end{cases}$$

Handwritten note: "the diff is h_{level+1}(k) and I know that"

- **SplitBucket()** splits bucket B_{next}

1. Redistribute the entries in B_{next} into B_{next+N_{level}} using h_{level+1}()
2. next = next + 1
3. if (next = N_{level}) then { level = level + 1; next = 0 } → you have split every bucket


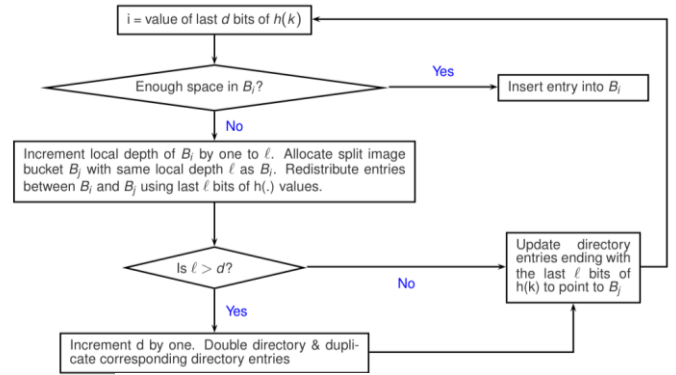
Linear Hashing Deletion

If last bucket is B_{level+next-1} is empty
 If next>0, next=-1
 If next=0 and level>0, next goes to last bucket of prev level then next=-1

Extensible Hashing Insert Overflow

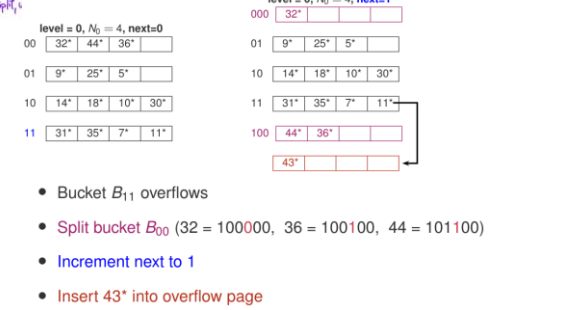
Each new directory entry (except for the entry for the split image) points to the same bucket as its corresponding entry
 Number of directory entries pointing to a bucket = $2^{d-\ell}$
 If split depth=global depth, double directory
 Else, split and increment local

Directory has 2^d entries
 d is called global depth of the hashed file
 Corresponding entries if their addresses differ only in the dth bit
 Each bucket maintains a local depth (denoted by $\ell \in [0, d]$)
 All entries in a bucket with local depth ℓ have the same last ℓ bits in h(.)



Extensible Hashing Deletion

If same local depth and differ only in ℓ th bit, merge
 If each pair of corresponding entries point to the same bucket, directory can be halved. $d=-1$



External Merge Sort

E.g. Sort 11-page using B=3

- 1. Read first 3 into B
- 2. Sort in main memory
- 3. Write sorted run to disk
- 4. Repeat 2 and 3
- 5. Merge each run, using B-1 for input, 1 for output

File size of N pages

Pass 0: Creation of sorted runs

- Read in and sort B pages at a time
- Number of sorted runs created = $\lceil N/B \rceil$
- Size of each sorted run = B pages

Pass i, $i \geq 1$: Merging of sorted runs

- Use B - 1 buffer pages for input & one buffer page for output
- Performs (B-1)-way merge

Analysis:

- N_0 = initial sorted runs = $\lceil N/B \rceil$
- +1 comes from initial pass
- Total number of passes = $\lceil \log_{B-1}(N_0) \rceil + 1$
- Total number of I/O = $2N(\lceil \log_{B-1}(N_0) \rceil + 1)$
- * Each pass reads N pages & writes N pages

Blocked I/O (optimize merge passes)

Read and write in units of buffer blocks of b
Given an allocation of B buffer pages for sorting,

Allocate one block (b pages) for output

- Remainder space can accommodate

$\lfloor (B-b)/b \rfloor$ blocks for input

- Thus, can merge at most $\lfloor (B-b)/b \rfloor$ sorted runs in each merge pass

Analysis:

N = number of pages in file to be sorted

B = number of available buffer pages

b = number of pages of each buffer block

N_0 = number of initial sorted runs = $\lceil N/B \rceil$

F = number of runs that can be merged at

each merge pass = $\lfloor B/b \rfloor - 1$

Num passes = $\lceil \log_F(N_0) \rceil + 1$

If $b=1$, then same as regular external merge

If b is good, then $\lceil \log_F(N_0) \rceil > 2N(\lceil \log_{B-1}(N_0) \rceil)$

With projection, size of record

divide by size of attributes

projected out

Page-nested needs ≥ 3 buffers

For block nested, check the

formula

When joining, size of record *2

cos R+S size

Always need 1 buffer for holding

input, output, and smaller table

Covering Index

An index I is a covering index for a query Q if all the attributes referenced in Q are part of the key or include column(s) of I
i.e. index includes all the attributes in Q.
esp if index is format1

- No RID lookup needed
- index-only plan

Conjunctive Normal Form (CNF) predicate



B+ Index matches non-disjunctive CNF predicate if at most one non-equality op (last attribute). Q: <=, =, then matches is at most <

Hash index matches p if all are equality

Primary Conjuncts

The conjuncts in p that I matches

Covered Conjunct

The conjuncts in p that I matches and appear in key or include columns of index

Notation

r relational algebra expression

$\lceil r \rceil$ number of tuples in output of r

$\lceil r \rceil$ number of pages in output of r

b_d num records on a page

b_l num entries on a page (leaf node)

F ave fanout of B+-tree index (i.e., number of pointers to child nodes)

h height of B+-tree inde: $h = \lceil \log_F(\lceil \frac{\lceil R \rceil}{b_l} \rceil) \rceil$

levels of internal nodes).

if format-2 index on table R

Cost of B+ tree index α Selection

C

$Cost_{internal} = \begin{cases} \lceil \log_F(\lceil \frac{\lceil R \rceil}{b_d} \rceil) \rceil & \text{if } I \text{ is a format-1 index,} \\ \lceil \log_F(\lceil \frac{\lceil R \rceil}{b_l} \rceil) \rceil & \text{otherwise.} \end{cases}$

$Cost_{leaf} = \begin{cases} \lceil \frac{\lceil \sigma_F(R) \rceil}{b_d} \rceil & \text{if } I \text{ is a format-1 index,} \\ \lceil \frac{\lceil \sigma_F(R) \rceil}{b_l} \rceil & \text{otherwise.} \end{cases}$

Hash-based (assume $B > \sqrt{f \times \lceil \pi_L^*(R) \rceil}$; i.e., no partition overflow)

Cost = $\lceil R \rceil + \lceil \pi_L^*(R) \rceil + \lceil \pi_L^*(R) \rceil$

Sort-based

- Output is sorted
- Good if there are many duplicates or if distribution of hashed values is non-uniform
- If $B > \sqrt{\lceil \pi_L^*(R) \rceil}$, then is desired

- * Number of initial sorted runs $N_0 = \lceil \frac{\lceil R \rceil}{B} \rceil \approx \sqrt{\lceil \pi_L^*(R) \rceil}$
- * Number of merging passes = $\log_{B-1}(N_0) \approx 1$
- * Sort-based approach requires 2 passes for sorting
- * Cost = $\lceil R \rceil + \lceil \pi_L^*(R) \rceil + \lceil \pi_L^*(R) \rceil$

Both hash-based & sort-based methods have same I/O cost

Hash Based Projection π_L

Consider $\pi_L(R)$

1 initialize an empty hash table T

2 for each tuple t in R do

3 apply hash function h on $\pi_L(t)$


4 let t be hashed to bucket Bi in T

5 if $(\pi_L(t) \text{ is not in } B_i)$ then

6 insert $\pi_L(t)$ into Bi

7 output all entries in T

Cost = $\lceil R \rceil$ if T fits in main memory



Partition

1 buffer input, (B - 1) buffers for output

Read R one page at a time into input buffer

For each tuple t in input buffer,

- project out unwanted attributes

- apply a hash function h on t' to distribute t' into one output buffer

- Flush to disk

Duplicate Elimination

For each partition R_i ,

- Initialize an in-memory hash table

- Read $\pi_L(R_i)$ one page at a time; for each tuple t read,

• Hash t into bucket Bj with hash function h'

(h' \neq h)

• Insert t into Bj if t \notin Bj

- Output tuples in hash table

Approach is effective if B is large relative to $\lceil R \rceil$

How large should B be?

- Assume that h distributes tuples in R uniformly
- Each R_i has $\lceil \frac{\lceil R \rceil}{B-1} \rceil$ pages
- Size of hash table for each $R_i = \lceil \frac{\lceil \pi_L^*(R) \rceil}{B-1} \rceil \times f$
 - * f = fudge factor > 1
- Therefore, to avoid partition overflow, $B > \lceil \frac{\lceil \pi_L^*(R) \rceil}{B-1} \rceil \times f$
 - * Approximately, $B > \sqrt{f \times \lceil \pi_L^*(R) \rceil}$

Analysis: Assume there's no partition overflow

- Cost of partitioning phase: $\lceil R \rceil + \lceil \pi_L^*(R) \rceil$
- Cost of duplicate elimination phase: $\lceil \pi_L^*(R) \rceil$
- Total cost = $\lceil R \rceil + 2\lceil \pi_L^*(R) \rceil$

No optimization sort

Cost Analysis

- Step 1:
 - Cost to scan records = $\lceil R \rceil$
 - Cost to output temporary result = $\lceil \pi_L^*(R) \rceil$
- Step 2:
 - Cost to sort records = $2\lceil \pi_L^*(R) \rceil (\log_m(N_0) + 1)$
 - N_0 = number of initial sorted runs, m = merge factor
- Step 3:
 - Cost to scan records = $\lceil \pi_L^*(R) \rceil$

Tuple Nested Loop Join $(\lceil R \rceil + \lceil R \rceil * \lceil S \rceil)$

for each tuple r \in R do

for each tuple s \in S do

if (r matches s) then

output (r,s) to result

Page Nested Loop Join $(\lceil R \rceil + \lceil R \rceil * \lceil S \rceil)$

for each page PR of R do

for each page PS of S do

for each tuple r \in PR do

for each tuple s \in PS do

if (r matches s) then

output (r,s) to result

Block Nested Loop Join

$\lceil R \rceil < \lceil S \rceil$. Allocate one page for S, one page for

output, remaining pages for R

while (scan of R is not done) do

read next (B - 2) pages of R into buffer

for each page P_s of S do

read P_s into buffer

for each r \in R in buffer, each s $\in P_s$ do

if (r matches s) then output (r,s)

Cost = $\lceil R \rceil + (\lceil \frac{\lceil R \rceil}{B-2} \rceil \times \lceil S \rceil)$

Index Nested Loop Join

Idea:

for each tuple r \in R do

use r to probe S's index to find matching tuples

Analysis:

- Let $R.A_i = S.B_j$ be the join condition
- Uniform distribution assumption:
 - each R-tuple joins with $\lceil \frac{\lceil S \rceil}{\lceil \pi_{B_j}(S) \rceil} \rceil$ number of S-tuples
- For a format-1 B+-tree index on S,

* I/O Cost = $\lceil R \rceil + \lceil R \rceil \times J$

* $J = \underbrace{\log_F(\lceil \frac{\lceil S \rceil}{b_d} \rceil)}_{\text{search index's internal nodes}} + \underbrace{\lceil \frac{\lceil S \rceil}{b_d \lceil \pi_{B_j}(S) \rceil} \rceil}_{\text{search index's leaf nodes}}$

(i) For 319-way merges, only 2 merging passes are needed. The first pass will produce 31250/319 = 98 sorted runs; these can then be merged in the next pass. Every page is read and written individually, at a cost of 16ms per read or write, in each of these two passes. The cost of these merging passes is therefore $2 \times 2 \times 16 \times 10000000 = 640,000,000ms$. (The formula can be read as number of passes times cost of read and write per page times number of pages in file.)

(ii) With 256-way merges, only 2 merging passes are needed. Every page in the file is read and written in each pass, but the effect of blocking is different on reads and writes. For reading, each page is read individually at a cost of 16ms. Thus, the cost of reads (over both passes) is $2 \times 16 \times 10000000 = 320000000ms$. For writing, pages are written out in blocks of 64 pages. The I/O cost per block is $10 + 5 + 1 \times 64 = 79ms$. The number of blocks written out per pass is $10000000/64 = 156250$, and the cost per pass is $156250 \times 79 = 12343750ms$.

of Computing, National University of Singapore

1. num merging passes
2. read per block x num blocks x passes
3. write per block x num blocks x passes
4. sum 2,3