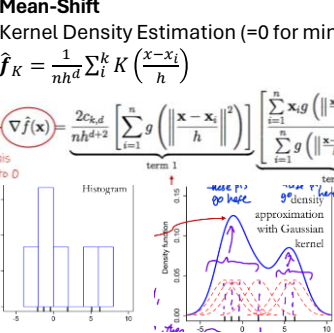


Kmeans
Params: l, or r, g, b, or x, y
- Setting k
- Sensitive to initial centres, outliers
- Assumes spherical clusters
- Assumes means are meaningful
SLIC

"#pixels, #superpixels: $n_{tp}; n_{sp}$
initial spacing of each superpixel $s = \sqrt{n_{tp}/n_{sp}}$
Features: $z = [r, g, b, x, y]$
Colour distance $d_c = ||(r_j - r_i, g_j - g_i, b_j - b_i)||$
Spatial Distance $d_s = ||(x_j - x_i, y_j - y_i)||$
Composite Distance $D = ||(\frac{d_c}{d_{cm}}, \frac{d_s}{d_{sm}})||$
 $\Rightarrow D = \sqrt{(d_c^2 + (d_s/s)^2 c^2)}$

Hyperparams: c(larger=constancy), n_{sp}
1. Initialise cluster centres to lowest gradient in 3x3 neighbourhood (avoid edges, noise), $label = -1, d(p) = \infty$
2. Assign, computing $D(p)$ between each ctr and pixel in $2s \times 2s$ neighbourhood.
3. Update cluster centres
4. Test convergence
5. Replace superpixels with average value

Mean-Shift
Kernel Density Estimation (=0 for minima)
 $\hat{f}_K = \frac{1}{nh^d} \sum_i K\left(\frac{x-x_i}{h}\right)$
 $\nabla \hat{f}(x) = \frac{2c_d}{nh^{d+2}} \left[\sum_{i=1}^n g\left(\left|\frac{x-x_i}{h}\right|\right) \left[\sum_{j=1}^n \mathbf{x}_j g\left(\left|\frac{\mathbf{x}-\mathbf{x}_i}{h}\right|\right) \right] - \mathbf{x} \right]$
Harris avoids even, eval
Vector x: $Ax = \lambda x$
let $A = \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix}$
To get λ , solve:
 $\det(A) = ad - bc = 0$.
Eval: solve $A \begin{bmatrix} x \\ y \end{bmatrix} = 0$
 $H = \begin{bmatrix} 0 & 0 \\ 0 & C \end{bmatrix}$
Vert edges: $I_y = 0$,
 $H = \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix}$
Cornerness score R
 $R = \min(\lambda_1, \lambda_2)$
 $\approx \det(H) - \text{trace}(H)$
 $= \lambda_1 \lambda_2 - \kappa(\lambda_1 + \lambda_2)^2$
 $(\text{trace}(A) = a + d)$



For each point, until convergence:
Initialise density window $x = x_j$
Compute mean shift vector (term 2)
Shift density window $x' = x + m(x), x = x'$
Hyperparam: h (bandwidth size)

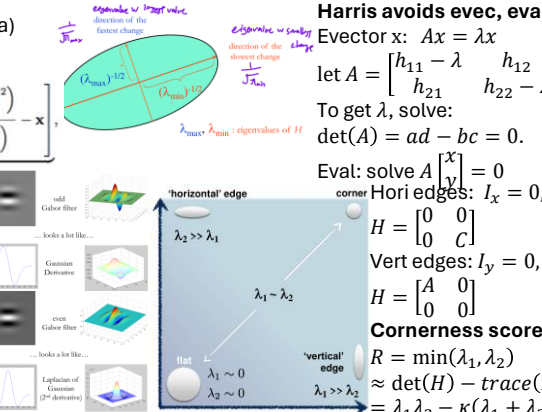
Optimisation
- all points within radius r have same window
- all points within search path have same window

Filter Banks
1. k-dim filter bank on m img: output: (m, y, x, k)
2. Kmeans and store cluster ctrs (Texton dict N-dim)
3. Each test img: filter w same bank for feat rep for each pixel. Assign to nearest cluster. Cluster ID is texton ID.


Histogram
Each pixel has closest ID from texton dict.
Hist can also be computed with in windows ctr'd per px neighbouring maxima of radius r

Gabor equation (sin=real, cos=even)
 $f_{mn} = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{m^2+y^2n^2}{2\sigma^2}\right) \sin\left(\frac{2\pi(\cos(\omega)m+\sin(\omega)n)}{\lambda} + \phi\right)$
 m, n : 2d coord of kernel. $m > n$, flat, $m < n$, skinny
 γ – aspect ratio of envelope coverage ($\gamma=1$ means circular, else elliptical)
 σ – rate of decay of exponential envelope (height of peaks)
 ω – orientation of sinusoid in 2D, influences orientation of detected feature (clockwise: 0 for vert, $\pi/2$ for hori)
 λ – wavelength of sinusoid; (wavelength is inversely proportional to frequency); (how spread-out peaks) usually set to be proportional to σ so that a constant number of cycles is visible)
 ϕ – phase of sinusoid; shift from center of kernel
Texture Boundary (dir edge detection)
Try all orientations of a disk Θ to find max hist difference

Textons are usually not (locally) precise for segmentation due to aggregation from window. Use spatially smoothed feature response instead.
Descriptors (save time complexity of shift_range)
Corner has significant change in all directions on shifting
Taylor expansion: $I(x+u, y+v) \approx I(x, y) + I_x u + I_y v$
SSD: $E(u, v) = \sum_{x,y \in W} [I(x+u, y+v) - I(x, y)]^2$
 $\approx \sum_{x,y \in W} [I_x u + I_y v]^2 = Au^2 + 2Buv + Cv^2$
 $= [u \ v] \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$, with $\Sigma A = I_x^2, \Sigma B = I_x I_y, \Sigma C = I_y^2$
Ellipse axes ori determined by the eigenvectors of H
Ellipse axes length determined by the eigenvalues of H




Harris Corner Detection
1. Compute gradient at each point
2. Convolve 3 channels w np.ones(window), calculate cornerness score $\det(H) - \kappa \text{tr}(H)$
3. Find points where $R > \text{threshold}$
4. NMS max points
Simple NMS
Search for max values and zero out surrounding window
Adaptive NMS
Pick corners where $R \gg \gg$ neighbouring maxima of radius r




Equivariance (transform img \rightarrow detection (locations) undergoes similar transformation)
Invariance: image transformed & detection (score) does not change
Harris Variance
Harris location is equivariant. response invariant to translation, rotation, additive intensity change. Variant to intensity scaling, img scaling, smoothing
LoG (blob detection)
 $\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$ roughly equals diff of gaussians
Gaussian: $f_{uv} = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{u^2+v^2}{2\sigma^2}\right)$
1st line: loc are equiv, 2nd line: descr are invar
Transformation: $I'(x, y) = I(x, y) + \Delta u$, where $\Delta u > 0$
custom, Harris, LoG, SIFT (all are invariant to intensity offset)
weighted intensity, weighted gradients, GIST, SIFT (all except GIST)
Transformation: $I'(x, y) = I(x, y) + \Delta u$, where $\Delta u > 0$
custom, Harris, LoG, SIFT (all are invariant to intensity offset)
weighted intensity, weighted gradients, GIST, SIFT (all except weighted intensity)
Transformation: $I'(x, y) = I(x, y)$ is a mirrored version of $I(x, y)$, i.e. flipped over the X axis.
custom, Harris, LoG, SIFT (all are invariant to flipping)
weighted intensity, weighted gradients, GIST, SIFT (GIST and SIFT are not mirror invariant)
Transformation: $I'(x, y) = I(y, x)$
custom, Harris, LoG, SIFT (all are invariant to transpose)
weighted intensity, weighted gradients, GIST, SIFT (GIST and SIFT descriptors are not invariant)
Transformation: $I'(x, y) = I(0.5 + x, y)$
custom, Harris, LoG, SIFT (none)
weighted intensity, weighted gradients, GIST, SIFT (none)
Transformation: $I'(x, y) = I(0.5 + x, 0.5 + y)$
custom, Harris, LoG, SIFT (all except custom)
weighted intensity, weighted gradients, GIST, SIFT (all except GIST)

Orientation Normalisation
In each window, take mode of grad bins
Multi-scale Oriented Patches (MOPS)
1. Take 40x40 window around detected keypoint (e.g. Harris corner)
2. Subsample every 5th pixel.
3. Rotate to horizontal.
4. Normalize intensity values within window $(x - \bar{x})/\sigma$
5. Wavelet transform on 8x8 patch to get a 64-dimensional descriptor
GIST (get rough spatial dist of img grad)
1. Divide image (patch) into 4 x 4 cells
2. Apply Gabor filters (directional edge detectors).
3. Compute filter response averages for each cell.
4. Size of descriptor is 4 x 4 x N, where N is size of the filter bank.

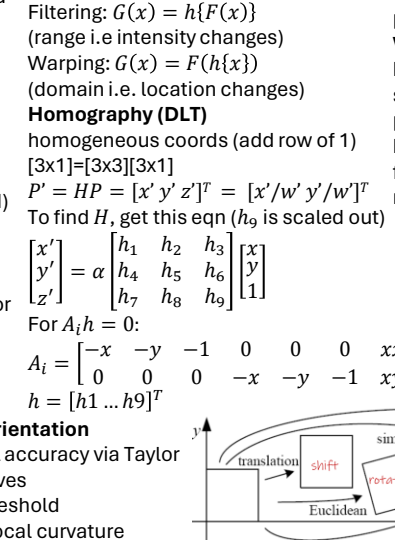
1. Multi-Scale Extrema (Blob) Detection
2. Localization + Scale/Orientation
• Refine kp loc to sub-pixel accuracy via Taylor expansion across the octaves
• Reject maxima below threshold
• Reject maxima w insuff local curvature



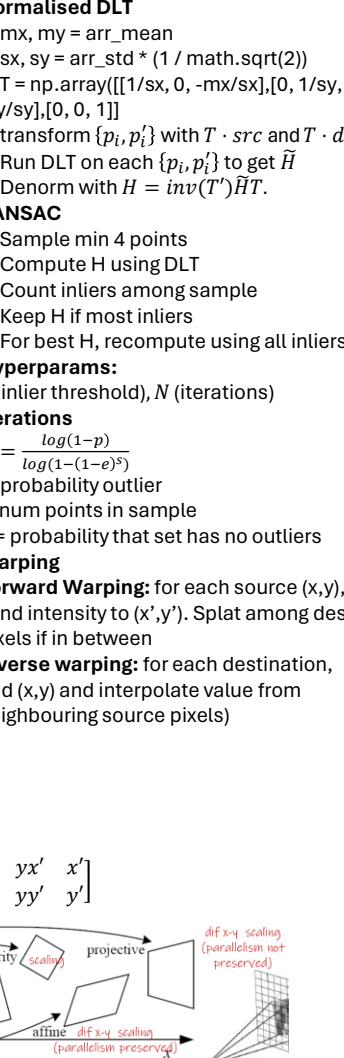
Solving for H
Given $\{p_i, p'_i\}$, for all p,
1. Create 2x9 matrix A_i
2. Concat into $2n \times 9$ matrix
3. Compute SVD of $A = U \Sigma V^T$
4. Store smallest singular v_i
5. Reshape to get H
DLT
- H has 8 dof, each kp correspondence is 2 eqn, so need 4 matches to solve
- Requires linear transformation (two images may not be projective)
- Sensitive to scaling (need norm)
- Sensitive to outliers
Normalised DLT
1. mx, my = arr_mean
2. sx, sy = arr_std * (1 / math.sqrt(2))
3. T = np.array([[1/sx, 0, -mx/sx], [0, 1/sy, -my/sy], [0, 0, 1]])
4. transform $\{p_i, p'_i\}$ with $T \cdot src$ and $T \cdot dst$
5. Run DLT on each $\{p_i, p'_i\}$ to get \tilde{H}
6. Denorm with $H = \text{inv}(T^T) \tilde{H} T$



Feature Matching
ratio $\text{dist} = ||f1 - f2|| / ||f1' - f2'||$
(best / next-best match)
Precision: TP/(TP+FP)
Recall: TP/(TP+FN)
Specificity: TN/(TN+FP)
ROC: TPR against FPR
Filtering: $G(x) = h\{F(x)\}$
(range i.e intensity changes)
Warping: $G(x) = F(h\{x\})$
(domain i.e. location changes)
Homography (DLT)
homogeneous coords (add row of 1)
 $[3 \times 1] = [3 \times 3] [3 \times 1]$
 $P' = HP = [x' \ y' \ z']^T = [x'/w' \ y'/w']^T$
To find H, get this eqn (h_9 is scaled out)
 $\begin{bmatrix} x' \\ y' \\ z' \end{bmatrix} = \alpha \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ h_7 & h_8 & h_9 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$
For $A_i h = 0$:
 $A_i = \begin{bmatrix} -x & -y & -1 & 0 & 0 & 0 & xx' & xy' & x' \\ 0 & 0 & 0 & -x & -y & -1 & xy' & yy' & y' \end{bmatrix}$
 $h = [h1 \dots h9]^T$




RANSAC
1. Sample min 4 points
2. Compute H using DLT
3. Count inliers among sample
4. Keep H if most inliers
5. For best H, recompute using all inliers
Hyperparams:
 δ (inlier threshold), N (iterations)
Iterations
 $N = \frac{\log(1-p)}{\log(1-(1-e^{\delta})^S)}$
e: probability outlier
s: num points in sample
p = probability that set has no outliers
Warping
Forward Warping: for each source (x,y), send intensity to (x', y') . Splat among dest pixels if in between
Inverse warping: for each destination, find (x,y) and interpolate value from neighbouring source pixels)



Optical Flow
Colour Constancy Assumption
 (Brightness constancy for intensity images)
 Implication: allows for pixel-to-pixel comparison (not image features)
Small Motion Assumption
 (pixels only move a little bit)
 Implication: linearization of the brightness constancy constraint

Optical Flow: (u, v)
Displacement: $(\delta x, \delta y) = (u \delta t, v \delta t)$
 $I(x + i \delta t, y + v \delta t, t + \delta t) = I(x, y, t)$
 doing Taylor expansion yields:


Brightness Constancy:
 $\frac{\partial I}{\partial x} \frac{\partial x}{\partial t} + \frac{\partial I}{\partial y} \frac{\partial y}{\partial t} + \frac{\partial I}{\partial t} = 0$
 $I_x u + I_y v + I_t = 0$
 $u = -\frac{I_y}{I_x} v - \frac{I_t}{I_x} \Rightarrow y = mx + b$



Lukas-Kanade Flow
 Constant Flow Assumption
 - all pixels in window have same constant flow


E.g. assume 5x5 patch has constant flow
 $\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$ (25 eqns)
 $Ax = -b \Rightarrow \hat{x} = \text{argmin} \|Ax - b\|^2$
 $\Rightarrow A^T A \hat{x} = A^T b$
 $\Rightarrow \begin{bmatrix} \sum_{p \in P} I_x I_x & \sum_{p \in P} I_x I_y \\ \sum_{p \in P} I_y I_x & \sum_{p \in P} I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum_{p \in P} I_x I_t \\ \sum_{p \in P} I_y I_t \end{bmatrix}$
 Solvable when $A^T A$ is invertible
 - λ_1 and λ_2 not too small
 - $A^T A$ well-conditioned
 - λ_1 / λ_2 not too large, where $\lambda_1 > \lambda_2$
 i.e. corners are good

If not small motion:



Aperture Problem
 When only line but no corners,
 I_x or $I_y = 0$ then unable to recover u or v
Coarse-to-Fine Estimation (pyramid)
 1. Downsample large motion
 2. Compute LK flow OF_0
 3. Upsample flow field
 4. Warp higher res image with LK flow
 5. Compute residual LK flow OF_1 using warped image
 6. Repeat w increasing higher-res img
 7. $OF = OF_0 + \dots \Delta OF_n$

Horn-Shunck Optical Flow
 Brightness constancy: every pixel,
 $\min [I_x u_{ij} + I_y v_{ij} + I_t]^2$
 Smooth flow field: $\min (u_{i,j} - u_{i+1,j})^2$
 Objective: $\min_{u,v} \sum_{i,j} \{E_s(i,j) + \lambda E_d(i,j)\}$



Ave of 4 neighbours:
HS Algorithm
 1. Precompute I_x, I_y, I_t
 2. Let $u = 0, v = 0$
 3. While not converged, (ave, frac, update eqns)

$$\bar{u}_{ij} = \frac{1}{4} \{u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1}\}$$

$$frac = \frac{I_x \bar{u}_{kl} + I_y \bar{v}_{kl} + I_t}{\lambda^{-1} + I_x^2 + I_y^2}$$

$$\hat{u}_{kl} = \bar{u}_{kl} - frac \times I_x$$

$$\hat{v}_{kl} = \bar{v}_{kl} - frac \times I_y$$

Hyperparams:
 λ small = emphasise smoothness
Summary:
LK Flow (constant flow), flow constant for all pixels, sparse, local method
 - Extracts visual features, tracks over multiple frames
 - Sparse motion fields
 - Suitable for large motion
HS Flow (smooth flow), varies per pixel, dense, global method
 - Based on spatio-temporal img brightness variations
 - Dense motion fields, sensitive to appearance variations

Ground Truth
 EndPointError = $\|\hat{f} - f_{GT}\|$
 Angular Error = $\arccos(\frac{\hat{f}^T f_{GT}}{|\hat{f}| |f_{GT}|})$
 where $f_{GT} = (u, v, 1)$, $\hat{f} = (\hat{u}, \hat{v}, 1)$

Apparent Motion
 - Projected Motion field may not be correct e.g. lighting changes
 - Constant Motion (LK) / smooth motion (HS) may not hold e.g. obj boundaries
 - Brightness constancy may not hold
 - Large motions, non-linear assumptions for iterative update

Tracking
 1. Template matching:
 Slow, combinatory, global solution
 2. Multi-scale template matching (norm-cc w smaller img to get peak, then upscale to get coarse location):
 Faster, combinatory, locally optimal
 3. Local refinement based on initial guess

Image Alignment
 Use prev frame as template for next frame
 Find warp parameter p for min SSD
 Guess p , assume small increment.
 Use Taylor series and solve using Least-Squares

Objective: $\min_p \sum_x [I(W(x; p)) - T(x)]^2$

$$= \min_{\Delta p} \sum_x [I(W(x; p_0 + \Delta p)) - T(x)]^2$$

$$= \min_{\Delta p} \sum_x \left[\nabla I \frac{\partial W}{\partial p} \Delta p - (T(x) - I(W(x; p_0))) \right]^2 \text{ equiv to } (Ax - b)$$

$$\Delta p = H^{-1} \sum_x \left[\nabla I \frac{\partial W}{\partial p} \right]^T [T(x) - I(W(x; p_0))]$$

$$H = \sum_x \left[\nabla I \frac{\partial W}{\partial p} \right]^T \left[\nabla I \frac{\partial W}{\partial p} \right]$$

x is coord, W is warp func (2x1) p_0 is warp param (2x3 for affine), ∇I is img grad (1x2), partial is Jacobian, Δp is incremental warp (6x1 for affine), $T(x)$ is template intensity

LK Alignment (fit template to img)
 1. Warp image $I(W(x; p))$
 2. Compute error image $T(x) - I(W(x; p_0))$
 3. Compute gradient $\nabla I(x')$
 4. Evaluate Jacobian $\frac{\partial W}{\partial p}$
 5. Compute Hessian (above)
 6. Compute Δp (above)
 7. Update p (above)

KLT Algorithm for Feature Tracking (track template across img)
 1. Find corners satisfying $\min(\lambda_1, \lambda_2) > \lambda$ for the Hessian
 2. Loop over corners:
 - Compute displacement to next frame using the LK align
 - Store displacement of each corner, update corner position
 - (opt) Add more corner pts every M frames using step 1
 3. Returns long trajectories for each corner point

Problems
 - Feature selection, efficiency
 - Accuracy eg rotation, shadow, drift
 - Occlusion (add, remove points)

Template-based Tracking:
 • Feature selection
 • RGB, gradient, histograms, deep features
 • Search for candidates
 • Limit search space based on prev results
 • Solve the mode-finding problem
 • Global vs local max
 • Update the target's template
 • Fixed vs update template

Metrics
 Accuracy
 - Intersection over union
 of bounding boxes (eg avg over frames)
 Robustness
 - Times that tracker fails (lose target)

Deep Learning Params
 - 200x200 pixels = 40k
 params per hidden unit

Normalised hist means scaling doesn't affect. But texon dict is affected

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & t_x \\ \sin(\theta) & \cos(\theta) & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x \cos(\theta) - y \sin(\theta) + t_x \\ x \sin(\theta) + y \cos(\theta) + t_y \end{bmatrix}$$

$$J = \begin{bmatrix} \frac{\partial x'}{\partial t_x} & \frac{\partial x'}{\partial t_y} & \frac{\partial x'}{\partial \theta} \\ \frac{\partial y'}{\partial t_x} & \frac{\partial y'}{\partial t_y} & \frac{\partial y'}{\partial \theta} \end{bmatrix} = \begin{bmatrix} 1 & 0 & -x \sin(\theta) - y \cos(\theta) \\ 0 & 1 & x \cos(\theta) - y \sin(\theta) \end{bmatrix}$$

DLT for circle

$$\begin{bmatrix} x & y & 1 \end{bmatrix} \begin{bmatrix} 2x_0 \\ 2y_0 \\ r^2 - x_0^2 - y_0^2 \end{bmatrix} = x^2 + y^2$$

Intier: $(y - y_i)^2 < \delta$

$$\text{eg } (r - \sqrt{(x - x_0)^2 + (y - y_0)^2})^2 < \delta$$

One brightness constancy eqn = 1 pixel = 1 equation
 SIFT is robust but not invariant to affine and projective transforms

MOUSE FILTER
 Minimise MSE of corr output
 $g = \text{argmin}_g \frac{1}{N} \sum_{i=1}^N (g \otimes x_i - y_i)^2 + \lambda \|g\|^2$
 $g = (X^T X + \lambda I)^{-1} X^T y$
 x is a $N \times hw$ matrix
 y is a $N \times hw$ output
 x_1 is first kernel, x_2, x_3 are subsequent kernels
Algorithm
 1. Initialise w several perturbed versions (robust)
 2. Optimise Kernel g in DFT
 3. Apply kernel g to frame I_k
 4. Find local maxima to locate target BB
 5. Update kernel w new BB

Jacobians (rmb to remove cols you don't need). $J = \square$ is just example

Translation eg $I'(x, y) = I(x + 3, y - 2)$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \end{bmatrix}$$

$$J = \begin{bmatrix} \frac{\partial x'}{\partial t_x} & \frac{\partial x'}{\partial t_y} \\ \frac{\partial y'}{\partial t_x} & \frac{\partial y'}{\partial t_y} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

Euclidean eg $I'(x, y) = I(-x, x)$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & t_x \\ \sin(\theta) & \cos(\theta) & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x \cos(\theta) - y \sin(\theta) + t_x \\ x \sin(\theta) + y \cos(\theta) + t_y \end{bmatrix}$$

$$J = \begin{bmatrix} \frac{\partial x'}{\partial t_x} & \frac{\partial x'}{\partial t_y} & \frac{\partial x'}{\partial \theta} \\ \frac{\partial y'}{\partial t_x} & \frac{\partial y'}{\partial t_y} & \frac{\partial y'}{\partial \theta} \end{bmatrix} = \begin{bmatrix} 1 & 0 & -x \sin(\theta) - y \cos(\theta) \\ 0 & 1 & x \cos(\theta) - y \sin(\theta) \end{bmatrix}$$

Similarity

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s \cos(\theta) & -s \sin(\theta) & t_x \\ s \sin(\theta) & s \cos(\theta) & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} sx \cos(\theta) - sy \sin(\theta) + t_x \\ sx \sin(\theta) + sy \cos(\theta) + t_y \end{bmatrix}$$

$$J = \begin{bmatrix} \frac{\partial x'}{\partial t_x} & \frac{\partial x'}{\partial t_y} & \frac{\partial x'}{\partial s} & \frac{\partial x'}{\partial \theta} \\ \frac{\partial y'}{\partial t_x} & \frac{\partial y'}{\partial t_y} & \frac{\partial y'}{\partial s} & \frac{\partial y'}{\partial \theta} \end{bmatrix} = \begin{bmatrix} 1 & 0 & x \cos(\theta) - y \sin(\theta) & -sx \sin(\theta) - sy \cos(\theta) \\ 0 & 1 & x \sin(\theta) + y \cos(\theta) & -sx \cos(\theta) + sy \sin(\theta) \end{bmatrix}$$

Affine eg $I'(x, y) = I(x + 0.3y + 3, 0.5x + y - 2)$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} p_1 & p_3 & t_x \\ p_2 & p_4 & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} p_1 x + p_3 y + t_x \\ p_2 x + p_4 y + t_y \end{bmatrix}$$

$$J = \begin{bmatrix} \frac{\partial x'}{\partial t_x} & \frac{\partial x'}{\partial t_y} & \frac{\partial x'}{\partial p_1} & \dots & \frac{\partial x'}{\partial p_4} \\ \frac{\partial y'}{\partial t_x} & \frac{\partial y'}{\partial t_y} & \frac{\partial y'}{\partial p_1} & \dots & \frac{\partial y'}{\partial p_4} \end{bmatrix} = \begin{bmatrix} 1 & 0 & x & 0 & y & 0 \\ 0 & 1 & 0 & x & 0 & y \end{bmatrix}$$

A, D are indicative of the aperture problem – they are all edge regions in which the resulting flow is perpendicular to the edge. B & C are correct; magnitude of the flow vectors seem high but this is likely due to scaling
 E – flow is correct – even though wall is not moving, the shadows create an “apparent” motion. This is not so visible due to the scaling in ground truth
 F – the swirl pattern of flow is incorrect, this is likely due to strong smoothing constraints from HS flow

Lukas-Kanade has sparse optical flow, HS has dense optical flow with smooth flow fields
 • Sparse-looking: B, C, F. Rule out B since flow vectors are still regular; so just C and F remaining. F is noisier and more flow vectors, so has the lower threshold for R;
 • Of remaining A, B, D, E for Horn-Shunck; A and D have more iterations – cannot recover the fine movements of the low-gradient areas such as the fabric background without running many iterations
 • Smaller lambda, smoother flow-fields so A is large lambda, B locally smoother than E.