

$f(x, y) = i(x, y) r(x, y)$
 $i(x, y)$: source illumination incident $0 \leq i(x, y) < \infty$
 $r(x, y)$: amount of illumination reflected by objects. $0 \leq r(x, y) \leq 1$ total absorption to total reflectance

Demosaicing Bayer Patterns

$[R, G, B]_B = [f(x+1, y+1), f(x+1, y), f(x, y)]$
 $[R, G, B]_{GB} = [f(x, y+1), f(x, y), f(x-1, y)]$
 $[R, G, B]_{GR} = [f(x+1, y), f(x, y), f(x, y-1)]$
 $[R, G, B]_R = [f(x, y), f(x-1, y), f(x-1, y-1)]$

Greyscale

$I = W_r R + W_g G + W_b B, \sum W_x = 1$

Normalized RGB:

$r/(R+G+B), g/(R+G+B), b/(R+G+B), I=(R+G+B)/3$

HSV


$$H = \begin{cases} \frac{G-B}{V-\min\{R, G, B\}} \cdot 60^\circ, & \text{if } V = R \text{ and } G \geq B; \\ \left(\frac{B-R}{V-\min\{R, G, B\}} + 2 \right) \cdot 60^\circ, & \text{if } G = V; \\ \left(\frac{R-G}{V-\min\{R, G, B\}} + 4 \right) \cdot 60^\circ, & \text{if } B = V; \\ \left(\frac{R-B}{V-\min\{R, G, B\}} + 5 \right) \cdot 60^\circ, & \text{if } V = R \text{ and } G < B \end{cases}$$

$$S = \frac{V - \min\{R, G, B\}}{V} \quad S \in [0, 1]$$

$$V = \max\{R, G, B\} \quad V \in [0, 255]$$

Gamma Mapping

$$x_{ij} = 255 \cdot \left(\frac{p_{ij}}{255} \right)^\gamma, \gamma > 0$$



Hist equalization

$x = \text{cdf}(x) \cdot \text{cap}$ eg: $0.3 \cdot 255$

Change intensity

Hist Stretching:

$x = (p - \min) \cdot \text{cap} / (\max - \min)$
eg: $(p - 50) \cdot 255 / (200 - 50)$

Stretch min and max

Gaussian Noise

Obs = Ideal + noise

$$p_{ij} = \hat{p}_{ij} + \eta$$

$$\eta \sim \mathcal{N}(\mu_n, \sigma_n)$$


mean $\mu = \frac{\sum_{i=1}^I \sum_{j=1}^J p_{ij}}{IJ}$

variance $\sigma^2 = \frac{\sum_{i=1}^I \sum_{j=1}^J (p_{ij} - \mu)^2}{IJ}$

resulting image pixels are zero-mean, unit variance

$$x_{ij} = \frac{p_{ij} - \mu}{\sigma}$$

Otsu's Method – Automated Thresholding



$$T^* = \underset{T}{\operatorname{argmin}} w_1(T) \cdot \sigma_1^2(T) + w_2(T) \cdot \sigma_2^2(T)$$

- optimal threshold T^* minimizes the sum of weighted variances of object and background.
- rationale:
 - correct threshold produces two narrow modes
 - incorrect threshold produces (at least one) wide mode
 - width of the mode is measured by the variance σ^2
 - weighting by number of pixels in each mode
- $\sigma_1^2(T), \sigma_2^2(T)$ - variance of pixels less than or equal to and greater than threshold respectively
- $w_1(T), w_2(T)$ - number of pixels less than or equal to and greater than threshold

we want to minimise sigma for each mode

Norm-CC

Conv/CC matches white

Norm-CC matches template

$$x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$$

$$x_{ij} = \frac{1}{|\mathbf{F}| |\mathbf{w}_{ij}|} \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$$

Sharpening:

stresses intensity peaks and differences with respect to the surroundings

Hori Sobel (Hori Grad=vert line)

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

Box Blur (Mean Filter)

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Sharpen

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

Outline

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 9 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Emboss

$$\begin{bmatrix} -2 & -1 & 0 \\ -1 & 1 & 1 \\ 0 & 1 & 2 \end{bmatrix}$$

Median filter

(denoise/delete pixels)

Vert Sobel (Vert grad=horiz line)

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Identity

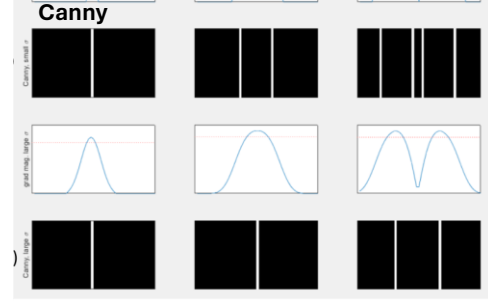
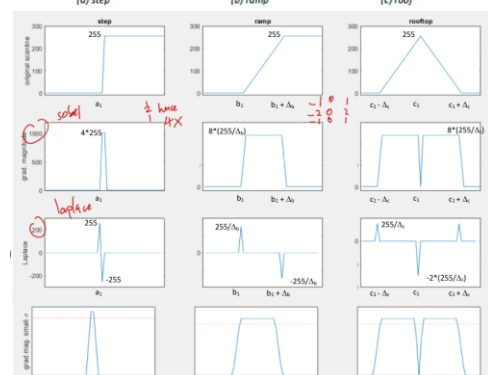

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Shift Right

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$

Laplace

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient is a vector that points in the direction of most rapid change in intensity.


$$\nabla f = \begin{bmatrix} 0 \\ 255 \end{bmatrix} \quad \nabla f = \begin{bmatrix} 0 \\ 255 \end{bmatrix}$$

The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

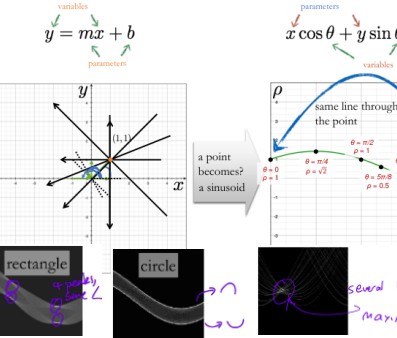


Normal Form

Line Detection with Hough Transform

Algorithm:

- Quantize Parameter Space (θ, ρ)
- Create Accumulator Array $A(\theta, \rho)$
- Set $A(\theta, \rho) = 0 \quad \forall \theta, \rho$
- For each image edge point (x_i, y_i)
For each element θ
Solve $\rho = x_i \cos \theta + y_i \sin \theta$
Increment $A(\theta, \rho) = A(\theta, \rho) + 1$
- Threshold, find local maxima in $A(\theta, \rho)$
- Detected line(s) given by $\rho^* = x \cos \theta^* + y \sin \theta^*$



Canny Edge Detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- Thin wide "ridges" down to single pixel width via non-maximum suppression
- Perform thresholding (hysteresis) & linking:
 - Mark pixels that pass high threshold as edge
 - Iteratively check and relabel pixels clearing low-threshold as edge if they are linked to an existing edge pixel

Hysteresis thresholding:

use a high threshold to start edge curves, and a low threshold to continue growing them. Iteratively mark adj pixels

- Find image gradients
 - G: gradient magnitude image (continuous range of values)
- Apply Non-Maximum Suppression to G
 - G_{NMS}: single pixel gradient image after NMS (continuous range of values)
- Apply double thresholding
 - Apply high threshold H to G_{NMS} to get E_H (binary output; pixels with values above or equal to threshold set to 1, pixels with values below set to 0)
 - Apply low threshold L to G_{NMS} to get E_L (binary output)
 - E_H = E_L ∪ E_L are weak pixels that passed the low threshold but failed the high threshold (binary output)
- Linking
 - E_{weak} = E_H ∪ E_L is the initial output edge image (binary)
 - E_{edge} = pixels of E_{weak} which are linked or connected to current E_{weak}
 - Update E_{weak} = E_{weak} ∪ E_{edge}; Remove E_{weak} from E_{weak} by setting to 0.
 - Repeat 4.1, 4.2, 4.3 while there are new pixels added to E_{weak} or some other stopping condition

Non-Maximum Suppression (NMS)

For each pixel, check if it is a local max along the gradient direction. If so, keep as edge, if not, discard (set to 0). May require checking interpolated values (e.g. pixel locations p and r). Alternative is to approximate with the closest pixel locations to p and r.


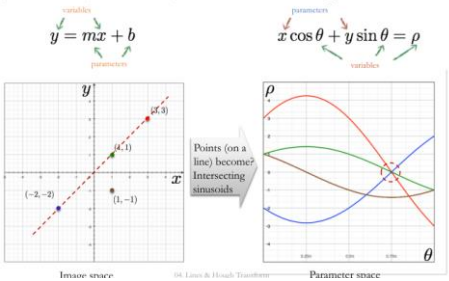


Image and (Normal) Parameter Space



There are two ways to write the same line: \cos we limit π and $-\pi$

Positive rho version:
 $x \cos \theta + y \sin \theta = \rho$

Recall:

$$\sin(\theta) = -\sin(\theta + \pi)$$

$$\cos(\theta) = -\cos(\theta + \pi)$$

Negative rho version:

$$x \cos(\theta + \pi) + y \sin(\theta + \pi) = -\rho$$

What if Radius is Unknown?


$(x-a)^2 + (y-b)^2 = r^2$ $(x-a)^2 + (y-b)^2 = r^2$

Radius is now a third parameter.

When considering the parameter space, this adds a third variable → makes the Hough space 3D.

- Augment accumulator array from $A(a, b)$ to $A(a, b, r)$
- In 2D parameter space, a point in image space corresponds to a circle (of radius r)
- In 3D parameter space, a point projected from image space also gets more complicated → cone


- Gradient information can save a lot of computation



Generalized Hough Transform

Offline Modeling


- Define a model shape by its boundary points p_i and reference point a .
- At each boundary point p_i , compute its displacement vector from a
 - Displacement vectors vote for potential a 's based on the position of p_i
- Store vectors in a table, indexed by p_i 's gradient orientation ϕ



Detection Procedure:

For each edge point:

- Use its gradient orientation ϕ to index into stored table
- Retrieve r vectors (mean/mode displacement vector for that ϕ) to vote for the reference point a



Voting

Minimize irrelevant tokens first, e.g. convert to edge image

Choose a good grid / discretization

Soft voting for neighbors (smoothing effect in accumulator array)

$$(\theta, \rho) \rightarrow 0.25 * (\theta, \rho - 1), 0.5 * (\theta, \rho), 0.25 * (\theta, \rho + 1)$$

Limit voting extent from each token

Use direction of edge to reduce amount of cast votes

Keep tags on votes to read back points contributing to local maxima

Feature Selection

Spatial

Intensity

Colour

k-Means Clustering

Basic idea: randomly initialize the k cluster centers, and iterate between assigning membership and computing cluster centers.

- Given K , randomly initialize the cluster centers, c_1, \dots, c_K
- Given cluster centers, determine points in each cluster
 - For each point p_i , find the closest c_j . Put p_i into cluster j
- Given points in each cluster, solve for c_j
 - Set c_j to be the mean of points in cluster j
- If c_j have changed (up to some threshold), repeat Steps 2, 3.

Pros

- Simple, fast to compute
- Converges to local minimum of within-cluster squared error

Cons/issues

- Setting k ?
- Sensitive to initial centers
- Sensitive to outliers
- Detects spherical clusters
- Assumes means can be computed (efficient, meaningful)

slope-intercept form: $y = mx + b$ (m =>slope, b =>y-intercept)

double-intercept form: $\frac{x}{a} + \frac{y}{b} = 1$ (a =>x-intercept, b =>y-intercept)

normal form: $x \cos \theta + y \sin \theta = p$

$y = -4x + 2$

$y + 4x = 2 \Rightarrow y/2 + x/(1/2) = 1$

$x \cos 14.07^\circ + y \sin 14.07^\circ = 2/(\sqrt{17})$

$\rho^2 = \frac{1}{\frac{1}{a^2} + \frac{1}{b^2}} = \frac{1}{(\frac{1}{2})^2 + 2^2} = \frac{4}{17}$, $\rho = \frac{2}{\sqrt{17}}$

Pros

All points are processed independently, so can cope with occlusion, gaps

Some robustness to noise: noise points unlikely to contribute consistently

Can detect multiple instances of a model in a single pass

Cons

Search time complexity increases exponentially with the # of model parameters

Non-target shapes can produce spurious peaks in parameter space

Quantization: can be tricky to pick a good grid size

SLIC

Definitions, Features & Distance Measure

n_p, n_{sp} → hyperparameter, normally determined with some other heuristic

Total number of pixels in image; number of superpixels we want

$s = [n_p / n_{sp}]^{1/2}$ Initial spacing of each superpixel

Features $z = [r, g, b, x, y]$

Colour Distance $d_c = [(r_j - r_i)^2 + (g_j - g_i)^2 + (b_j - b_i)^2]^{1/2}$

Spatial Distance $d_s = [(x_j - x_i)^2 + (y_j - y_i)^2]^{1/2}$

Composite Distance $D = \left[\left(\frac{d_c}{d_{cm}} \right)^2 + \left(\frac{d_s}{d_{sm}} \right)^2 \right]^{1/2} \Rightarrow D = [d_c^2 + \left(\frac{d_s}{s} \right)^2]^{1/2}$

normalisation: set as max expected values of d_c and d_s

$d_{cm} = s = [n_p / n_{sp}]^{1/2}$

scaling factors: since color is [0, 255] range, image can be very big → set to same scale

if you use HSV, the use of H & S can difficult to make to converge faster

grey-scale version

L2 distance measures

c becomes nearly hyperparameter

see FPA for derivation

Consider as hyperparameter weighting relative importance between colour similarity vs. spatial proximity, s

Initialization

- Initialize the algorithm: Compute the initial superpixel cluster centers,

$$m_i = [r_i, g_i, b_i, x_i, y_i]^T, i = 1, 2, \dots, n_{sp}$$

by sampling the image at regular grid steps, s . Move the cluster centers to the lowest gradient position in a 3×3 neighborhood. For each pixel location, p , in the image, set a label $L(p) = -1$ and a distance $d(p) = \infty$.

Different from k-means; rather than initialize randomly, initialize cluster centers on a grid.

Do not initialize cluster center on an edge or noise point (these have some more extreme values); convergence is faster if we can already initialize on a value common to the surrounding pixels, (approximate as lowest gradient position).





05. Image Segmentation

Mean Shift Algorithm

Main Idea: find **modes** or local density maxima in the feature space

For each data point:

- Define a window around it, compute the centroid
- Shift the center of the window to the centroid
- Repeat until the centroid stops moving (convergence).



Mean-Shift Procedure

For each point x_j , where $j = 1 \dots n$

- Initialize density window x : $x = x_j$
- Compute mean shift vector m :
Mean shift vector is a function of the bandwidth h and kernel G . These are expressed as subscripts to keep notation tidy.
- Shift density window and update: $x' = x + m(x)$
 $x = x'$
- Iterate steps 2 and 3 until **convergence**.

Segmentation

- Find features (color, gradients, texture, etc) to represent pixels.
- Initialize density windows at individual feature points.
- Perform mean shift for each point until convergence.
- Merge pixels whose feature points that end up near the same mode or "peak" into the same segment.

(Need to choose bandwidth/window size)

Pros

General algorithm for mode-finding (we apply it towards segmentation)

No prior assumptions on cluster shape (spherical, elliptical, etc.)

One parameter (window size / bandwidth h , which has a physical meaning)

Finds variable number of modes (vs. pre-specified k in k-means)

Robust to outliers

Cons

Output depends on h and selecting h is non-trivial

Computationally expensive and slow to run!

Scales poorly with feature space dimension

Optimisations

Points along search path/end point are of same cluster